



TURNING TWEETS INTO KNOWLEDGE

An Introduction to Text Analytics



Twitter

- Twitter is a social networking and communication website founded in 2006
- Users share and send messages that can be no longer than 140 characters long
- One of the Top 10 most-visited sites on the internet
- Initial Public Offering in 2013
- Valuation ~\$31 billion



Impact of Twitter

- Use by protestors across the world
- Natural disaster notification, tracking of diseases
- Celebrities, politicians, and companies connect with fans and customers
- Everyone is watching!



2013 AP Twitter Hack



- The Associated Press is a major news agency that distributes news stories to other news agencies
- In April 2013 someone tweeted the above message from the main AP verified Twitter account
- S&P500 stock index fell 1% in seconds, but the White House rapidly clarified

Understanding People

- Many companies maintain online presences
- Managing public perception in age of instant communication essential
- Reacting to changing sentiment, identifying offensive posts, determining topics of interest...
- How can we use analytics to address this?



Using Text as Data

- Until now, our data has typically been

- Structured
- Numerical
- Categorical



hannah @lawlorff

1h

MY ELECTRIC HAS WENT OUT AND A GIANT SPIDER IS COMING 4 ME And mY ONLY SOURCE OF LIGHT IS THE FLASHLIGHT ON MY PHONE GOD BLESS @Apple

Expand

- Tweets are

- Loosely structured
- Textual
- Poor spelling, non-traditional grammar
- Multilingual



matt @clairvoyant

2h

WHYCANT I GO BACK TO IOS6 ITS NOT THAT BIG A DEAL @Apple
I LIKE YOUR OLD OPERATING SYSTEM BETTER

Expand

Text Analytics



- We have discussed why people care about textual data, but how do we handle it?
- Humans can't keep up with Internet-scale volumes of data
 - ~500 million tweets per day!
- Even at a small scale, the cost and time required may be prohibitive

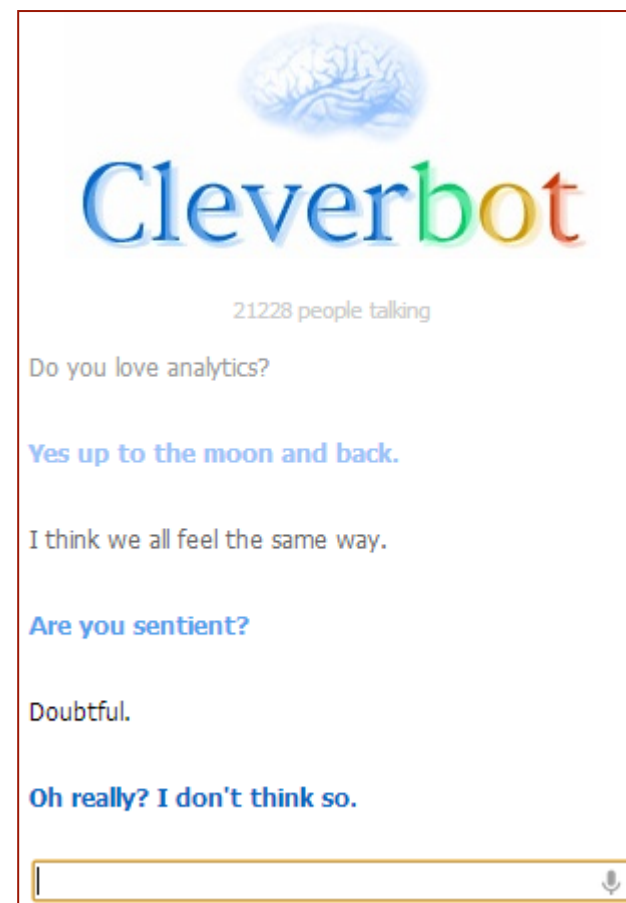
How Can Computers Help?

- Computers need to understand text
- This field is called **Natural Language Processing**
- The goal is to understand and derive meaning from human language
- In 1950, Alan Turing proposes a test of machine intelligence: passes if it can take part in a real-time conversation and cannot be distinguished from a human



History of Natural Language Processing

- Some progress: “chatterbots” like ELIZA
- Initial focus on understanding grammar
- Focus shifting now towards statistical, machine learning techniques that learn from large bodies of text
- Modern “artificial intelligences”: Apple’s Siri and Google Now



Why is it Hard?

- Computers need to understand text
- Ambiguity:
 - “I put my **bag** in the **car**. **It** is **large and blue**”
 - “**It**” = **bag**? “**It**” = **car**?
- Context:
 - Homonyms, metaphors
 - Sarcasm
- In this lecture, we’ll see how we can build analytics models using text as our data

Sentiment Mining - Apple

- **Apple** is a computer company known for its laptops, phones, tablets, and personal media players
- Large numbers of fans, large number of “haters”
- Apple wants to monitor how people feel about them over time, and how people receive new announcements.
- **Challenge: Can we correctly classify tweets as being negative, positive, or neither about Apple?**



Creating the Dataset



- Twitter data is publically available
 - Scrape website, or
 - Use special interface for programmers (API)
 - Sender of tweet may be useful, but we will ignore
- Need to construct the outcome variable for tweets
 - Thousands of tweets
 - Two people may disagree over the correct classification
 - One option is to use **Amazon Mechanical Turk**

Amazon Mechanical Turk




- Break tasks down into small components and distribute online
- People can sign up to perform the tasks for a fee
 - Pay workers, e.g. \$0.02 per classified tweet
 - Amazon MTurk serves as a broker, takes small cut
- Many tasks require human intelligence, but may be time consuming or require building otherwise unneeded capacity

Our Human Intelligence Task

- Actual question we used:

Judge the sentiment expressed by the following item toward the software company "Apple"

- Workers could pick from
 - Strongly Negative (-2)
 - Negative (-1)
 - Neutral (0)
 - Positive (+1)
 - Strongly Positive (+2)



- Five workers labeled each tweet

Our Human Intelligence Task

- For each tweet, we take the average of the five scores.
 - “LOVE U @APPLE” (1.8)
 - “@apple @twitter Happy Programmers' Day folks!” (0.4)
 - “So disappointed in @Apple Sold me a Macbook Air that WONT run my apps. So I have to drive hours to return it. They wont let me ship it.” (-1.4)
- We have labels, but how do we build independent variables from the text of a tweet?

A Bag of Words

- Fully understanding text is difficult
- Simpler approach:

Count the number of times each words appears

- “This course is great. I would recommend this course to my friends.”

THIS	COURSE	GREAT	...	WOULD	FRIENDS
2	2	1	...	1	1

A Simple but Effective Approach



- One feature for each word - a simple approach, but effective
- Used as a baseline in text analytics projects and natural language processing
- Not the whole story though - preprocessing can dramatically improve performance!

Cleaning Up Irregularities

- Text data often has many inconsistencies that will cause algorithms trouble
- Computers are very literal by default – Apple, APPLE, and ApPLe will all be counted separately.
- Change all words to either lower-case or upper-case

Apple	APPLE	ApPLe		apple
apple	apple	apple		3

Cleaning Up Irregularities

- Punctuation also causes problems – basic approach is to remove everything that isn't a,b,...,z
- Sometimes punctuation **is** meaningful
 - Twitter: **@apple** is a message to Apple, **#apple** is about Apple
 - Web addresses: [www.website.com/somepage.html](#)
- Should tailor approach to the specific problem

@Apple	APPLE!	--apple--		apple
apple	apple	apple		3

Removing Unhelpful Terms

- Many words are frequently used but are only meaningful in a sentence - “**stop words**”
 - Examples: *the, is, at, which...*
 - Unlikely to improve machine learning prediction quality
 - Remove to reduce size of data
- Two words at a time?
 - “**The Who**” → “ ”
 - “Take **That**” → “Take”



Stemming

- Do we need to draw a distinction between the following words?

argue argued argues arguing

- Could all be represented by a common **stem**, **argu**
- Algorithmic process of performing this reduction is called **stemming**
- Many ways to approach the problem

Stemming

- Could build a **database of words** and their stems
 - **Pro:** handles exceptions
 - **Con:** won't handle new words, bad for the Internet!
- Can write a **rule-based** algorithm
 - e.g. if word ends in “ed”, “ing”, or “ly”, remove it
 - **Pro:** handles new/unknown words well
 - **Con:** many exceptions, misses words like **child** and **children** (but would get other plurals: **dog** and **dogs**)

Stemming

- The second option is widely popular
 - “**Porter Stemmer**” by Martin Porter in 1980, still used!
 - Stemmers have been written for many languages
- Other options include machine learning (train algorithms to recognize the roots of words) and combinations of the above

Real example from data:

*“by far the best **customer** care **service** I have ever **received**”*

⇒ *“by far the best **custom** care **servic** I have ever **receiv**”*

Sentiment Analysis Today



- Over 7,000 research articles have been written on this topic
- Hundreds of start-ups are developing sentiment analysis solutions
- Many websites perform real-time analysis of tweets
 - “tweetfeel” shows trends given any term
 - “The Stock Sonar” shows sentiment and stock prices

Text Analytics in General



- Selecting the specific features that are relevant in the application
- Applying problem specific knowledge can get better results
 - Meaning of symbols
 - Features like number of words

The Analytics Edge



- Analytical sentiment analysis can replace more labor-intensive methods like polling
- Text analytics can deal with the massive amounts of unstructured data being generated on the internet
- Computers are becoming more and more capable of interacting with humans and performing human tasks
- In the next lecture, we'll discuss IBM Watson, an impressive feat in the area of Text Analytics